

# Java Meets Rust: The Depths of Cross-Language Integration



# Introduction



**Konstantin Grechishchev**

 [linkedin.com/in/kgrech](https://www.linkedin.com/in/kgrech)

kmgrechis@gmail.com



## Experience

- Master of computer science, 2015
- Cisco Systems Inc



**WhatsApp**

## WhatsApp Software Engineer

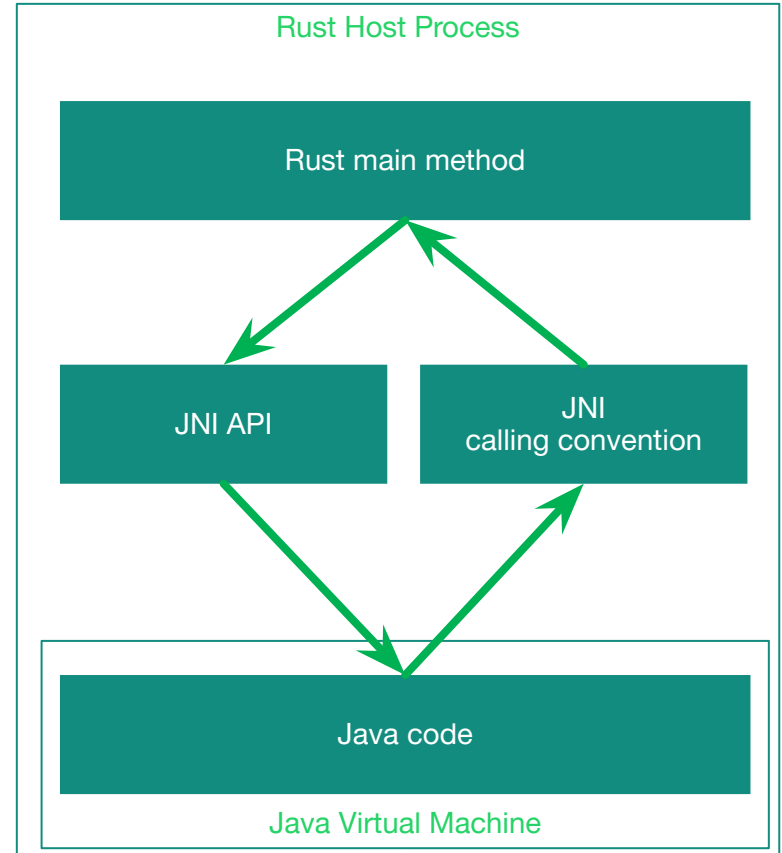
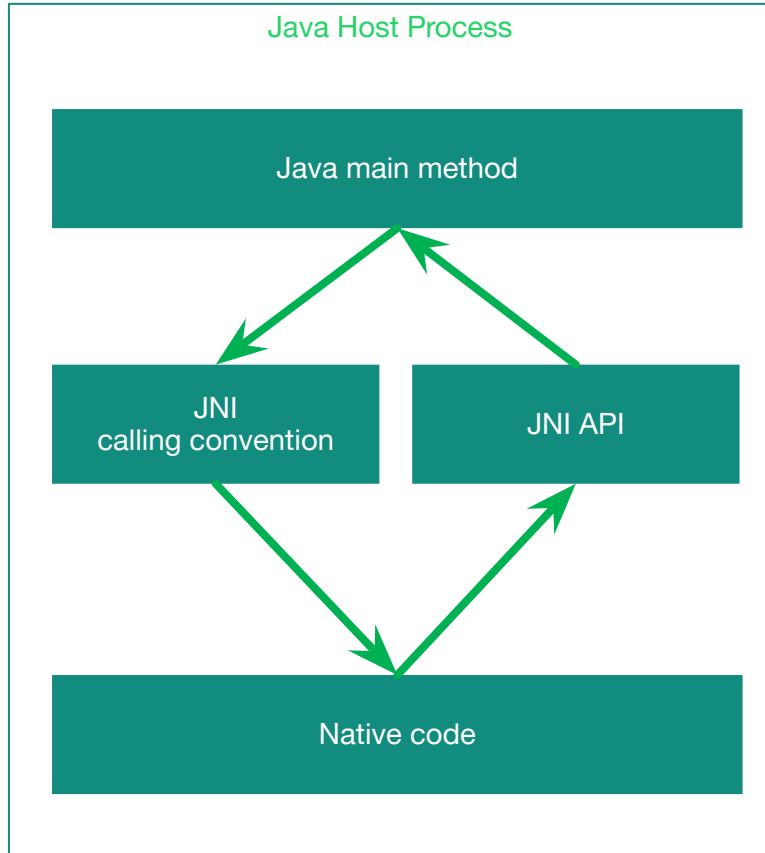
- Since 2020

# Do you really need rust and Java in the same process?

- Maybe you don't?
- Can you keep them as separate processes?
- Ways of interprocess communication
  - Rest
  - GRPC/Thrift/etc
  - TCP/UWP/Web Sockets
  - Files
  - Database
  - Message queues
  - Operation system IPC mechanisms
- Same process, but different architecture?
  - CRUX?



# Rust and Java in the same process





# How to pass complex data?

- Use JNI APIs to create or traverse java objects ([jni-rs](#),
  - Freedom to define interface according to your preference
  - Could be broken by minifier and obfuscator
  - Requires aligned changes on java and rust
  - A lot of work!
- Generate FFI interface ([UniFFI](#) - kotlin)
  - Restricted to the capabilities of the framework
  - Not so many options available (write yours?)
- Serialize data on one side and deserialize on other (protobuf or even json!)
  - Thin JNI layer
  - Write once and only modify the schema later
  - Slightly less efficient (profile before jurgung!)
  - Boilerplate code to serialize and deserialize



# Meet Toy JNI!

- Define protobuf schema
- Generate bindings for both java and rust
- Pass arrays of bytes over JNI interface
- Implement wrappers to make it more convenient!
- Done!
- Simply extend schema when required



```
syntax = "proto3";

package toy_jni;
option java_package = "com.github.kgrech.toyjni.proto";
option java_multiple_files = true;

message Request {
    string message = 1;
    uint64 response_delay = 2;
}

message Response {

    message Success {
        string message = 1;
    }

    message Error {
        string error_message = 1;
    }

    optional Success success = 1;
    optional Error error = 2;
}
```

# JNI Calling Convention


- [Oracle Java Native Interface Specification. Chapter 2: Design Overview](#)

- Java:

```
package pkg;
class Cls {
    native double f(int i, String s);
    static {
        // Consider compiling all-rust dependencies in single library
        System.loadLibrary("libname");
    }
}
```

- C:

```
#include <jni.h>
JNIEXPORT jdouble JNICALL Java_pkg_Cls_f (
    JNIEnv *env,          /* interface pointer */
    jobject obj,          /* "this" pointer */
    jint i,               /* argument #1 */
    jstring s)            /* argument #2 */
{
    /* Obtain a C-copy of the Java string */
    const char *str = (*env)->GetStringUTFChars(env, s, 0);
    /* process the string */
    /* Now we are done with str */
    (*env)->ReleaseStringUTFChars(env, s, str);
    return ...
}
```



# Java calling Rust (jni crate)

- Declare the native method

```
private native byte[] nativeCall(byte[] request);
```

- Define the implementation in rust

```
#[no_mangle]
pub extern "system" fn Java_com_github_kgrech_toyjni_sync_JNIBridge_nativeCall<'a>(
    env: JNIEnv<'a>,
    j_object: jobject<'a>,
    request: JByteArray<'a>,
) -> JByteArray<'a> {
    handle_error(env, |env| {
        let request = Request::decode_from_java(env, &request)?;
        // Do something with request
        let response = if request.message == "Hello, Rust!" {
            Response::success("Hello, Java!".into())
        } else {
            Response::error(format!("Unable to respond to '{}'", request.message))
        };
        response.encode_to_java(env)
    })
}
```



# Decoding data efficiently!

- [Oracle Java Documentation. Chapter 4: JNI Functions](#)

- GetByteArrayRegion - **copies** a region of a primitive array into a buffer.
- GetByteArrayElements - returns the body (**maybe a copy**) of the primitive array, valid until ReleaseByteArrayElements().
- GetPrimitiveArrayCritical/ReleasePrimitiveArrayCritical - similar to GetByteArrayElements, but asks JVM to avoid copy. There are restrictions on what could be done between Get and Release calls.

- JNI crate provides nice wrappers:

```
fn decode_from_java(env: &mut JNIEnv, j_byte_array: &JByteArray) -> anyhow::Result<Self>
{
    unsafe {
        let guard = env
            .get_array_elements_critical(j_byte_array, ReleaseMode::NoCopyBack)
            .with_context(|| "Failed to access byte array")?;
        let slice: &[i8] = &*guard;
        let u8_slice: &[u8] =
            std::slice::from_raw_parts(slice.as_ptr() as *const u8, slice.len());
        Self::decode(u8_slice).with_context(|| "Unable to decode request")
    }
}
```

# Encoding the data

- [Oracle Java Documentation. Chapter 4: JNI Functions](#)

- NewByteArray - construct a new primitive array object
- SetByteArrayRegion - **copies** a region of a primitive array from a buffer.

- JNI crate provides a nice wrapper for both:

```
fn encode_to_java<'a>(&self, env: &JNIEnv<'a>) -> anyhow::Result<JByteArray<'a>> {  
    let protobuf_bytes: Vec<u8> = self.encode_to_vec();  
    env.byte_array_from_slice(protobuf_bytes.as_slice())  
        .with_context(|| "Error converting result to java byte array")  
}
```

# Calling Java Back

- Let's define a function to be called from rust:

```
public byte[] onJniCallback(byte[] request) { // Called from native code
    Request deserializedRequest = Request.parseFrom(request);
    Response response = handleCallbak(deserializedRequest);
    return response.toByteArray();
}
```


- [Oracle Java Documentation. Chapter 4: JNI Functions](#)

- Call<type>Method Routines, Call<type>MethodA Routines, Call<type>MethodV Routines used to call a Java instance method from a native method

```
let data = response.encode_to_java(env)?;
let response = env
    .call_method(
        j_object,
        "onJniCallback",
        "([B)[B", // Method signature
        &[JValue::from(&data)] ,
    )
    .with_context(|| "Error calling onJniCallback method"?;
```

# JNIEnv lifetimes

```
pub extern "system" fn Java_com_github_kgrech_toyjni_sync_JNIBridge_nativeCall<'a>(  
    env: JNIEnv<'a>,  
    j_object: JObject<'a>,  
    request: JByteArray<'a>,  
) -> JByteArray<'a>
```



Note the lifetimes!

- [Oracle Java Documentation. Chapter 5: The Invocation API](#)
- JNIEnv is a pointer to native method interface (main set of the APIs). **Valid only in the current thread**
- Could be obtained:
  - As a parameter to any call made from Java
  - From JavaVM \*jvm using AttachCurrentThread/AttachCurrentThreadAsDaemon JNI APIs
- Lifetime ensures JVM is not terminated and the object is not garbage collected
- This won't compile:

```
pub extern "system" fn nativeCall<'a>(env: JNIEnv<'a>, j_object: JObject<'a>) {  
    std::thread::spawn(move || {  
        callback(env, j_object); // env and j_object are neither Send nor 'static!  
    });  
}
```

# How do we pass the Java object to another thread?

- [Oracle Java Documentation. Chapter 4: JNI Functions](#)

- NewGlobalRef - Creates a new global reference, must be explicitly disposed of by DeleteGlobalRef().
- GetJavaVM - Returns the Java VM interface (used in the Invocation API) associated with the current thread

```
let global_ref: GlobalRef = env.new_global_ref(j_object)
let jvm: JavaVM = env.get_java_vm()
```

**Have no lifetime.  
Implement Send.**

- JavaVM is a JNI interface pointer (JVM lifecycle and threads API) and denotes a Java VM.
- Could be obtained:
  - If you start the JVM by calling JNI\_CreateJavaVM()
  - From JNIEnv env using GetJavaVM method





# Storing Java Object inside Rust Struct

- Why?
  - To keep it for a call done in background
- Can't store `j_object: JObject<'a>`
  - It has a lifetime. Because Java will GC collect it!
- Use [NewGlobalRef](#) API

```
let global_ref: GlobalRef = env.new_global_ref(j_object)?;
```

- Global references must be explicitly disposed of by calling the [DeleteGlobalRef](#) method.
  - Rust wrapper calls it called on drop
- Can use `as_obj` to get `JObject` back for API call when needed

```
let j_object: JObject = global_ref.as_obj();
```

- Do not forget to attach the thread!

# Rust Part of the Callback (1)

- Protect JObject from being collected garbage collected and get JVM pointer:

```
pub struct CallbackContext {
    jvm: JavaVM,
    global_ref: GlobalRef,
}

impl CallbackContext {
    pub fn new(env: &JNIEnv, j_object: &JObject) -> anyhow::Result<Self> {
        let global_ref = env
            .new_global_ref(j_object)
            .with_context(|| "Unable to create global reference for java object repo")?;
        let jvm: JavaVM = env
            .get_java_vm()
            .with_context(|| "Unable to JVM reference")?;
        Ok(Self { jvm, global_ref })
    }
}
```

## Rust Part of the Callback (2)

- CallbackContext implements Send trait, so we can pass it to another thread safely

```
pub extern "system" fn nativeCall<'a>(env: JNIEnv<'a>, j_object: JObject<'a>) {  
    let context = CallbackContext::new(env, &j_object)?;  
    std::thread::spawn(move || {  
        std::thread::sleep(  
            Duration::from_millis(request.response_delay)  
        );  
        let response = context.callback(Request::message(  
            "Hello, Java!".into()  
        ))  
        .expect("Callback error!");  
        println!("Response from callback: {response:?}")  
    });  
}
```

# Beware: Attaching a thread

- [Oracle Java Documentation. Chapter 5: The Invocation API](#)
- AttachCurrentThread API attaches the current thread to a Java VM.
  - You can think about it as “set required thread-local variables”
  - Returns a JNI interface pointer in the JNIEnv argument.
  - Attaching a thread that is already attached is a no-op, can't attach one thread to two VMs.
  - You must attach thread in order to call any JNI API (rust enforces it via JNIEnv lifetime)



## Rust Part of the Callback (3)

- Attach the thread and make the call

```
impl CallbackContext {  
    pub fn callback(&self, response: Request) -> anyhow::Result<Response> {  
        // Attach current thread, guard derefs to JEnv  
        let mut guard = self.jvm.attach_current_thread();  
        let data = response.encode_to_java(&guard)?;  
        // Call java back  
        let response = guard // Derefs to &JNIEnv  
            .call_method(  
                self.global_ref.as_obj(),  
                "onJniCallback",  
                "([B)[B",  
                &[JValue::from(&data)],  
            );  
  
        // Convert response to JObject  
        let j_byte_array: JByteArray = response.l().into();  
        Response::decode_from_java(&mut guard, &j_byte_array)  
    }  
}
```



# Lets go async?



# Async method to call

- We want to call `async_handle_request` method from Java:

```
async fn async_handle_request(request: Request) -> Response {
    if request.response_delay > 0 {
        // Emulate async IO
        tokio::time::sleep(Duration::from_millis(request.response_delay)).await;
    }
    let response = if request.message == "Hello, Rust!" {
        Response::success("Hello, Java!".into())
    } else {
        Response::error(format!("Unable to respond to '{}'", request.message))
    };
    response
}
```

# Can't await in synchronous methods

```
#[no_mangle]
pub extern "system" fn Java_com_github_kgrech_toyjni_async_JNIBridge_nativeCall<'a>(
    env: JNIEnv<'a>,
    j_object: JObject<'a>,
    request: JByteArray<'a>,
) -> JByteArray<'a> {
    handle_error(env, |env| unsafe {
        let request = Request::decode_from_java(env, &request)?;
        let response = async_handle_request(request).await;
        response.encode_to_java(env)
    })
}
```

- Compile error: “only allowed inside `async` functions and blocks”.
- We need to access an async runtime!

# Storing the runtime

- Option 1: make it global and store in OnceCell.
- Option 2: Associate the runtime pointer with java object

- Java

```
public class JNIBridge {  
    private final long runtime;  
    public JNIBridge() { // Can't have native constructor  
        this.runtime = init();  
    }  
    private native long init();  
}
```

- Rust

```
#[no_mangle]  
pub extern "system" fn Java_com_github_kgrech_toyjni_async_JNIBridge_init<'a>(  
    _env: JNIEnv<'a>, _j_object: JObject<'a>,  
) -> jlong {  
    let runtime = Arc::new(Runtime::new().expect("Unable to init runtime"));  
    Arc::into_raw(runtime) as jlong // Convert arc to pointer, forget it and cast to long  
}
```

# Accessing the runtime

- Access the field back:
- Get<Type>Field APIs are the way to go

/// Converts the value of the field into Arc **without** incrementing the Arc's ref counter.

```
pub unsafe fn get_field<T>(  
    env: &mut JNIEnv,  
    j_object: &JObject,  
    field_name: &str,  
) -> anyhow::Result<Arc<T>> {  
    let field_value: JValueOwned = env  
        .get_field(j_object, field_name, Primitive::Long.to_string())  
        .with_context(|| format!("Can't get long value of the '{field_name}' field"))?;  
    let j_long_value: jlong = field_value  
        .j()  
        .with_context(|| format!("Wrong type of the '{field_name}' field.")).?;  
    let ptr = j_long_value as *const c_void;  
    Ok(Arc::from_raw(ptr.cast()))  
}
```



# Cleanup

- Java

```
public class JNIBridge implements AutoCloseable {  
    @Override  
    public native void close();  
}
```

- Rust

```
const RUNTIME_FIELD: &str = "runtime";  
  
#[no_mangle]  
pub extern "system" fn Java_com_github_kgrech_toyjni_async_JNIBridge_close<'a>(  
    mut env: JNIEnv<'a>,  
    j_object: JObject<'a>,  
) {  
    unsafe {  
        let _runtime = get_field::<Runtime>(&mut env, &j_object, RUNTIME_FIELD);  
        // drop(_runtime); is called implicitly  
    }  
}
```

# Accessing the runtime

- Borrow instead

```
/// Converts the value of the field into Arc and increments the Arc's ref counter
pub unsafe fn borrow_field<T>(
    env: &mut JNIEnv,
    j_object: &JObject,
    field_name: &str,
) -> anyhow::Result<Arc<T>> {
    let handle: Arc<T> = get_field(env, j_object, field_name)?;
    let clone = handle.clone(); /// Increment ref-counter
    std::mem::forget(handle); /// Forget the Arc again!
    Ok(clone)
}
```

- Now we can call it this way:

```
let runtime: Arc<Runtime> = borrow_field(env, &j_object, RUNTIME_FIELD)? ;
```

# Beware: passing the raw pointer to Java

- **Unsafe**
  - Ensure there is no code flow allowing user API input to change long value
  - Consider `OnceCell<GlobalContext>` for sensitive cases
- **Optimizations**
  - Java would complain that the field is not used (annotate with `@Keep?`)
- **Access by name**
  - Ensure the field is not renamed during obfuscation
- **Concurrency**
  - Use `Box::into_raw` and `Box::from_raw` if you guarantee exclusive thread safe access from Java
  - Use `Arc<T>` if the value is read only
  - Otherwise, consider `Arc<Mutex<T>>` or `Arc<RwLock<T>>`



# Blocking method

- Simple block-on:

```
#[no_mangle]
pub extern "system" fn Java_com_github_kgrech_toyjni_async_JNIBridge_nativeCall<'a>(
    env: JNIEnv<'a>,
    j_object: jobject<'a>,
    request: JByteArray<'a>,
) -> JByteArray<'a> {
    handle_error(env, |env| unsafe {
        let request = Request::decode_from_java(env, &request)?;
        let runtime: Arc<Runtime> = borrow_field(env, &j_object, RUNTIME_FIELD)?;
        let response = runtime.block_on(async_handle_request(request));
        response.encode_to_java(env)
    })
}
```

- Java thread is blocked
  - Bad idea to block main thread in Android or coroutine dispatcher threads
  - Could end up spawning too many threads

# Use Java abstraction for delayed computation

- Java has [Future](#) and [CompletableFuture](#) class
- Return CompletableFuture from the native code

```
private native CompletableFuture<byte[]> nativeCall(byte[] request);
```

- Create instance using [FindClass](#) and [NewObject](#) APIs. `env.new_object` is a wrapper for both

```
const FUTURE_CLASS: &str = "java/util/concurrent/CompletableFuture";
```

```
const FUTURE_CONSTRUCTOR: &str = "()V";
```

```
pub struct ToyJniFuture {  
    /// Pointer to java future object. GlobalRef garbage collection  
    java_future: GlobalRef,  
}  
  
impl ToyJniFuture {  
    pub fn new(env: &mut JNIEnv) -> anyhow::Result<Self> {  
        let java_future_obj = env.new_object(FUTURE_CLASS, FUTURE_CONSTRUCTOR, &[])?;  
        let java_future = env.new_global_ref(java_future_obj)?;  
        Ok(Self { java_future })  
    }  
}
```



# Beware: Classloaders

- Java Classloaders are associated with a caller class, not with a thread!
- If you create a thread in the rust code, it has no stack frames from java.
- The "system" class loader to be used
- Your app might have custom associated classloader
- Options:
  - Do your FindClass lookups in the method called by JVM
  - Cache a reference to the ClassLoader object somewhere, then manually call the loadClass method
- See more: [Why didn't FindClass find my class?](#)



# Completing the future

- When ready, call complete method:

```
const COMPLETE_METHOD_NAME: &str = "complete";
const COMPLETE_METHOD_SIGNATURE: &str = "(Ljava/lang/Object;)Z";
impl ToyJniFuture {
    pub fn complete(&self, env: &mut JNIEnv, object: &JObject) -> anyhow::Result<bool> {
        let response: JValueOwned = env
            .call_method(
                self.java_future.as_obj(),
                COMPLETE_METHOD_NAME,
                COMPLETE_METHOD_SIGNATURE,
                &[JValue::from(object)],
            )
            .with_context(|| "Error calling complete"?);
        response
            .z()
            .with_context(|| "Method returned non-bool value")
    }
}
```

- Note that we need &mut JNIEnv and JObject

# Callback context

```
pub struct CallbackContext { // Send + Sync + `static
    jvm: JavaVM, // Keep the JVM pointer
    future: ToyJniFuture, // And the future
}

impl CallbackContext {
    pub fn callback(&self, response: Response) -> anyhow::Result<()> {
        let mut guard = self // Attach current thread to the JVM
            .jvm
            .attach_current_thread()
            .with_context(|| "Unable to attach current thread to the JVM"?;
        if let Some(success) = response.success {
            let data = success.encode_to_java(&guard)?;
            self.future.complete(&mut guard, &data)?;
        } else { todo!() }
        Ok(())
    }
}
```


# Implement native call

- Implement `Java_com_github_kgrech_toyjni_async_JNIBridge_nativeCall`:

```
#[no_mangle]
pub extern "system" fn Java_com_github_kgrech_toyjni_async_JNIBridge_nativeCall<'a>(...)
-> JObject<'a> {
    handle_error_exceptionally(env, |env| unsafe {
        let request = Request::decode_from_java(env, &request)?;
        let runtime: Arc<Runtime> = borrow_field(env, &j_object, RUNTIME_FIELD)?;
        let future = ToyJniFuture::new(env)?;
        let local_ref = future.local_reference(env)?;
        // If required, load required classes here and include them to CallbackContext
        let callback_context = CallbackContext::new(env, future)?;
        runtime.spawn(async move {
            let response = async_handle_request(request).await;
            callback_context
                .callback(response)
                .expect("Unrecoverable callback error");
        });
        Ok(local_ref)
    })
    .unwrap_or_default()
}
```

# Beware: JMV calls in the async runtime thread

- Requires attachment and detachment of the async executor thread to JVM
- Could cause performance bottlenecks
- Consider:
  - Attach dedicated thread to VM
  - Dispatch computation results over the sync channel to this thread
  - Loop over result queue and call complete methods



# Error handling

- Implement `Java_com_github_kgrech_toyjni_async_JNIBridge_nativeCall`:

```
#[no_mangle]
pub extern "system" fn Java_com_github_kgrech_toyjni_async_JNIBridge_nativeCall<'a>(...)
-> JObject<'a> {
    handle_error_exceptionally(env, |env| unsafe {
        let request = Request::decode_from_java(env, &request)?;
        let runtime: Arc<Runtime> = borrow_field(env, &j_object, RUNTIME_FIELD)?;
        let future = ToyJniFuture::new(env)?;
        let local_ref = future.local_reference(env)?;
        // If required, load required classes here and include them to CallbackContext
        let callback_context = CallbackContext::new(env, future)?;
        runtime.spawn(async move {
            let response = async_handle_request(request).await;
            callback_context
                .callback(response)
                .expect("Unrecoverable callback error");
        });
        Ok(local_ref)
    })
    .unwrap_or_default()
}
```

# Throwing Exceptions

- Calling `unwrap` causes a panic in rust and likely to crash the JVM :(
- [Throw and ThrowNew](#) API are the way to go

```
const TOY_JNI_EXCEPTION: &str = "com/github/kgrech/toyjni/ToyJNIException";
pub fn throw_exception(env: &mut JNIEnv, error: impl Display) {
    // Check if there is an existing Java exception
    match env.exception_check() {
        Ok(true) => {
            // If there is an existing exception, return here to re-throw it
        }
        Ok(false) => {
            let result = env.throw_new(TOY_JNI_EXCEPTION, error.to_string());
            if let Err(err) = result {
                env.fatal_error(format!("Error throwing {TOY_JNI_EXCEPTION}: {err}"))
            }
        }
        Err(err) => env.fatal_error(format!("Exception check failed: {err}")),
    }
}
```

# Passing exception to the future (1)

- Convert error the the exception and pass it
- Modify CallbackContext::callback:

```
impl CallbackContext {  
  
    pub fn callback(&self, response: Response) -> anyhow::Result<()> {  
        // Existing code to attach current thread to the JVM  
  
        if let Some(success) = response.success {  
            ...  
        } else if let Some(error) = response.error {  
            let exception = create_exception(&mut guard, error.error_message)?;  
            self.future.complete_exceptionally(&mut guard, &exception)?;  
        }  
        Ok(())  
    }  
}
```



# Creating Exceptions

- It is possible to create the instances of exceptions same as other Java objects
- [FindClass and NewObject](#) APIs could be used. `env.new_object` is a wrapper for both

```
const TOY_JNI_EXCEPTION: &str = "com/github/kgrech/toyjni/ToyJNIException";
const TOY_JNI_EXCEPTION_CONSTRUCTOR: &str = "(Ljava/lang/String;)V";

pub fn create_exception<'a>(env: &mut JNIEnv<'a>, error: impl Display) ->
    anyhow::Result<JThrowable<'a>> {
    let java_sting = env.new_string(error.to_string())
        .with_context(|| "Error converting error string to Java string")?;
    let j_object = env.new_object(
        TOY_JNI_EXCEPTION,
        TOY_JNI_EXCEPTION_CONSTRUCTOR,
        &[JValue::from(&java_sting)],
    )
        .with_context(|| format!("Error creating {TOY_JNI_EXCEPTION}"))?;
    Ok(j_object.into())
}
```

## Passing exception to the future (2)

- Call the `completeExceptionally(Throwable ex)` method
- The same pattern as earlier:

```
const COMPLETE_EXCEPTIONALLY_METHOD_NAME: &str = "completeExceptionally";
const COMPLETE_EXCEPTIONALLY_METHOD_SIGNATURE: &str = "(Ljava/lang/Throwable;)Z";

impl ToyJniFuture {
    pub fn complete_exceptionally(&self, env: &mut JNIEnv, object: &JThrowable) ->
        anyhow::Result<bool> {
        let response = env
            .call_method(
                self.java_future.as_obj(),
                COMPLETE_EXCEPTIONALLY_METHOD_NAME,
                COMPLETE_EXCEPTIONALLY_METHOD_SIGNATURE,
                &[JValue::from(object)],
            )
            .with_context(|| "Error calling completeExceptionally"?;
        response.z().with_context(|| "method returned non-bool value")
    }
}
```

## How it looks in java?

[illegible]

# Thank you! Q&A

- [Oracle JNI Specification](#)
- [The Java™ Native Interface](#)
- <https://www.linkedin.com/in/kgrech/>
- kmgrechis@gmail.com



<https://github.com/kgrech/toy-jni>